

Zerocash, Bitcoin, and Transparent Computational Integrity

Eli Ben-Sasson, Technion

Based on joint works with Iddo Ben-Tov, Alessandro Chiesa, Michael Forbes, Ariel Gabizon, Daniel Genkin, Matan Hamilis, Ynon Horesh, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Nick Spooner, Eran Tromer, Madars Virza

January 2017

Overview

- Computational integrity and privacy (CIP) — motivation
- Importance of transparency
- A pair of new transparent CI(P) systems

Overview

- Computational integrity and privacy (CIP) — motivation
- Importance of transparency
- A pair of new transparent CI(P) systems

Brief bio:

- Professor at Technion — Israel Institute of Technology
- started w/ theoretical CS related to efficient proofs (“moon math”)

Overview

- Computational integrity and privacy (CIP) — motivation
- Importance of transparency
- A pair of new transparent CI(P) systems

Brief bio:

- Professor at Technion — Israel Institute of Technology
- started w/ theoretical CS related to efficient proofs (“moon math”)
- noticed (circa 2008) proofs efficient enough to reduce to practice
- traveled to Bitcoin 2013 to “show and tell” ...

Overview

- Computational integrity and privacy (CIP) — motivation
- Importance of transparency
- A pair of new transparent CI(P) systems


Brief bio:

- Professor at Technion — Israel Institute of Technology
- started w/ theoretical CS related to efficient proofs (“moon math”)
- noticed (circa 2008) proofs efficient enough to reduce to practice
- traveled to Bitcoin 2013 to “show and tell” ...
- ... Zerocash co-author, ZCash co-founder & scientist

Overview

- Computational integrity and privacy (CIP) — motivation
- Importance of transparency
- A pair of new transparent CI(P) systems

Brief bio:

- Professor at Technion — Israel Institute of Technology
- started w/ theoretical CS related to efficient proofs (“moon math”)
- noticed (circa 2008) proofs efficient enough to reduce to practice
- traveled to Bitcoin 2013 to “show and tell” ...
- ... Zerocash co-author, ZCash co-founder & scientist
- also interested in *CRowd-based IInteractive Curation (CROINC)*
- early childhood development tracker: Baby.CROINC.org  **baby Croinc**

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),
- may wish to preserve privacy of parts of the data (*privacy problem*)

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),
- may wish to preserve privacy of parts of the data (*privacy problem*)
- examples: tax reporting, proof of (fractional) reserve, shielded payments, . . .

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),
- may wish to preserve privacy of parts of the data (*privacy problem*)
- examples: tax reporting, proof of (fractional) reserve, shielded payments, . . .
- the CIP problem: guaranteeing correct output reported, without compromising privacy

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),
- may wish to preserve privacy of parts of the data (*privacy problem*)
- examples: tax reporting, proof of (fractional) reserve, shielded payments, . . .
- the CIP problem: guaranteeing correct output reported, without compromising privacy
- Zero knowledge proofs/arguments [GMR88] use randomness, interaction and cryptography to solve both CI and P in an astonishingly efficient way;

Computational Integrity and Privacy (CIP)

A party executing program P on mix of public/private data . . .

- may have *incentive* to *misreport* the output (*integrity problem*),
- may wish to preserve privacy of parts of the data (*privacy problem*)
- examples: tax reporting, proof of (fractional) reserve, shielded payments, . . .
- the CIP problem: guaranteeing correct output reported, without compromising privacy
- Zero knowledge proofs/arguments [GMR88] use randomness, interaction and cryptography to solve both CI and P in an astonishingly efficient way;
- protocols solving CIP are also known as protocols for *checking* [BFL91], *certifying* [M94], *delegating* [GKR08], and *verifying* [GGP10], computations

IP, ZK, MIP, PCP, NIZK, CS [circa 1990]

Definition (Computational Integrity (CI))

is the language of quadruples $(M, \mathcal{T}, x_{\text{in}}, x_{\text{out}})$ such that nondeterministic machine M , on input x_{in} reaches output x_{out} after \mathcal{T} cycles, \mathcal{T} in binary.

IP, ZK, MIP, PCP, NIZK, CS [circa 1990]

Definition (Computational Integrity (CI))

is the language of quadruples $(M, \mathcal{T}, x_{\text{in}}, x_{\text{out}})$ such that nondeterministic machine M , on input x_{in} reaches output x_{out} after \mathcal{T} cycles, \mathcal{T} in binary.

Lemma

CI is NEXP-complete

IP, ZK, MIP, PCP, NIZK, CS [circa 1990]

Definition (Computational Integrity (CI))

is the language of quadruples $(M, \mathcal{T}, x_{\text{in}}, x_{\text{out}})$ such that nondeterministic machine M , on input x_{in} reaches output x_{out} after \mathcal{T} cycles, \mathcal{T} in binary.

Lemma

CI is NEXP-complete

Definition (proof system)

An proof system S for L is a pair $S = (V, P)$ satisfying

- **efficiency** V is randomized polynomial time; P unbounded
- **completeness** $x \in L \Rightarrow \Pr[V(x) \leftrightarrow P(x) \rightsquigarrow \text{accept}] = 1$
- **soundness** $x \notin L \Rightarrow \Pr[V(x) \leftrightarrow P(x) \rightsquigarrow \text{accept}] \leq 1/2$

IP, ZK, MIP, PCP, NIZK, CS [circa 1990]

Definition (Computational Integrity (CI))

is the language of quadruples $(M, \mathcal{T}, x_{\text{in}}, x_{\text{out}})$ such that nondeterministic machine M , on input x_{in} reaches output x_{out} after \mathcal{T} cycles, \mathcal{T} in binary.

Lemma

CI is NEXP-complete

Definition (argument system)

An **argument** system S for L is a pair $S = (V, P)$ satisfying

- **efficiency** V is randomized polynomial time; P is similarly bounded
- **completeness** $x \in L \Rightarrow \Pr[V(x) \leftrightarrow P(x) \rightsquigarrow \text{accept}] = 1$
- **soundness** $x \notin L \Rightarrow \Pr[V(x) \leftrightarrow P(x) \rightsquigarrow \text{accept}] \leq 1/2$

IP, ZK, MIP, PCP, NIZK, CS [circa 1990]

Definition (Computational Integrity (CI))

is the language of quadruples $(M, \mathcal{T}, x_{\text{in}}, x_{\text{out}})$ such that nondeterministic machine M , on input x_{in} reaches output x_{out} after \mathcal{T} cycles, \mathcal{T} in binary.

Lemma

CI is NEXP-complete

Theorem ([BM88, GMR88, BFL88, BFL91, BGKW88, FLS90, BFLS91, AS92, ALMSS92, K92, M94])

CI has an argument system $S = (V, P)$ that is

- **noninteractive:** Prover sends a single message (requires setup/RO)
- **succinct:** Verifier run-time $\text{poly}(n, \log \mathcal{T})$; this bounds proof length
- **transparent:** Setup+verifier queries are public random coins
- **zero knowledge:** proof preserves privacy of nondeterministic witness

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
 - ▶ TPs have served societies for millenia
 - ▶ TPs *want* to stay such, so are not incentivized to pay for crypto CIP
 - ▶ Cryptographic CIP seems computationally costly compared to TP model (encryption suffices)

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
 - ▶ TPs have served societies for millenia
 - ▶ TPs *want* to stay such, so are not incentivized to pay for crypto CIP
 - ▶ Cryptographic CIP seems computationally costly compared to TP model (encryption suffices)
 - ▶ but considering the costs of *manual* CIP (audits, legislation, regulation), cryptographic CIP is cheap!

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
- Enter Bitcoin !
 - ▶ decentralized, “In Crypto we trust”
 - ▶ huge incentive to compromise integrity ($1\text{BTC} > 1,000\$$ (1/1/2017))
 - ▶ privacy crucial for fungibility and business-adoption

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
- Enter Bitcoin !
- Zerocash, ZCash
 - ▶ Bitcoin Conference (5/2013), offered using CIP to improve Bitcoin
 - ★ blockchain compression
 - ★ proof of reserve ("I own 1,000 BTC")
 - ★ improved privacy

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
- Enter Bitcoin !
- Zerocash, ZCash
 - ▶ Bitcoin Conference (5/2013), offered using CIP to improve Bitcoin
 - ★ blockchain compression
 - ★ proof of reserve ("I own 1,000 BTC")
 - ★ improved privacy
 - ▶ Zerocoin (5/2013): RSA-accumulator decentralized mix [MGGR13]

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
- Enter Bitcoin !
- Zerocash, ZCash
 - ▶ Bitcoin Conference (5/2013), offered using CIP to improve Bitcoin
 - ★ blockchain compression
 - ★ proof of reserve ("I own 1,000 BTC")
 - ★ improved privacy
 - ▶ Zerocoin (5/2013): RSA-accumulator decentralized mix [MGGR13]
 - ▶ Zerocash [B,Chiesa,Garman,Green,Miers,Tromer,Virza 14]
 - ★ first Decentralized Anonymous Payment (DAP) system
 - ★ hides payer, payee and payment amount
 - ★ uses KOE-based zkSNARKs [GGPR13,BCGTV13]

Who needs cryptographic CIP?

- Trusted parties (TP)? banks, Google, Facebook, Visa, PayPal, ...
- Enter Bitcoin !
- Zerocash, ZCash

Overview

- Computational integrity and privacy (CIP) — motivation ✓
- Importance of transparency
- A pair of new transparent CI(P) systems

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time
 - ▶ any code you write, I can prove it was executed correctly

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time
 - ▶ any code you write, I can prove it was executed correctly
- succinctness — short proofs, fast verification
 - ▶ compression of computation (like checking the whole blockchain), without compromising integrity and privacy
 - ▶ “heavy-weight proving” done only once, only by prover

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time
 - ▶ any code you write, I can prove it was executed correctly
- succinctness — short proofs, fast verification
 - ▶ compression of computation (like checking the whole blockchain), without compromising integrity and privacy
 - ▶ “heavy-weight proving” done only once, only by prover
- Zerocash/ZCash uses (zkSNARKs) that achieve the above
 - ▶ based on bilinear pairings [G10,GGP10,L12] and Quadratic Arithmetic Programs (QAP) [GGPR13]
 - ▶ these zkSNARKs require *non-transparent* setup phase
 - ▶ if setup compromised, leaks a forgery-trapdoor

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time
 - ▶ any code you write, I can prove it was executed correctly¹
- succinctness — short proofs, fast verification
 - ▶ compression of computation (like checking the whole blockchain), without compromising integrity and privacy
 - ▶ “heavy-weight proving” done only once, only by prover
- Zerocash/ZCash uses (zkSNARKs) that achieve the above
 - ▶ based on bilinear pairings [G10,GGP10,L12] and Quadratic Arithmetic Programs (QAP) [GGPR13]
 - ▶ these zkSNARKs require *non-transparent* setup phase
 - ▶ if setup compromised, leaks a forgery-trapdoor

¹fine print: assuming I don't know the trapdoor

CIP — desirable features

- universality — arbitrary programs (NP/NEXP complete)
- scalability — efficient prover running-time
 - ▶ any code you write, I can prove it was executed correctly¹
- succinctness — short proofs, fast verification
 - ▶ compression of computation (like checking the whole blockchain), without compromising integrity and privacy
 - ▶ “heavy-weight proving” done only once, only by prover
- Zerocash/ZCash uses (zkSNARKs) that achieve the above
 - ▶ based on bilinear pairings [G10,GGP10,L12] and Quadratic Arithmetic Programs (QAP) [GGPR13]
 - ▶ these zkSNARKs require *non-transparent* setup phase
 - ▶ if setup compromised, leaks a forgery-trapdoor
- **Definition:** A CIP system is *transparent* if setup and all verifier queries are public random coins (Arthur-Merlin protocol)

¹fine print: assuming I don't know the trapdoor

Transparency important for

- Ongoing public trust in integrity of the system
 - ▶ even one trapdoor leak could completely ruin integrity
 - ▶ increased value \Rightarrow increased incentive to attack/corrupt
 - ▶ what if powerful entity/agency asks for the trapdoor?

Transparency important for

- Ongoing public trust in integrity of the system
 - ▶ even one trapdoor leak could completely ruin integrity
 - ▶ increased value \Rightarrow increased incentive to attack/corrupt
 - ▶ what if powerful entity/agency asks for the trapdoor?
- Collaborative creation of CIP software
 - ▶ crypto-currencies use decentralized code development
 - ▶ who generates keys for a non-transparent CI?
 - ▶ with code proliferation, should you trust the setup?

Transparency important for

- Ongoing public trust in integrity of the system
 - ▶ even one trapdoor leak could completely ruin integrity
 - ▶ increased value \Rightarrow increased incentive to attack/corrupt
 - ▶ what if powerful entity/agency asks for the trapdoor?
- Collaborative creation of CIP software
 - ▶ crypto-currencies use decentralized code development
 - ▶ who generates keys for a non-transparent CI?
 - ▶ with code proliferation, should you trust the setup?
- Transparent auditing of central private registries
 - ▶ registries maintained by governments have huge impact on citizen rights and liabilities
 - ▶ many registries contain private data, so privacy prevents public auditing
 - ▶ cryptographic CIP can enhance trust in registry management
 - ▶ public trust demands transparent CIP

Overview

- Computational integrity and privacy (CIP) — motivation ✓
- Importance of transparency ✓
- A pair of new transparent CI(P) systems

A pair of novel transparent CIP

① SCI (Scalable Computational Integrity)

- ▶ Joint work with Iddo Ben-Tov, Alessandro Chiesa, Ariel Gabizon, Daniel Genkin, Matan Hamilis, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Eran Tromer and Madars Virza
- ▶ To appear in Eurocrypt 2017
- ▶ universal, succinct, scalable, transparent, post-quantum secure

② SCIP (Scalable Computational Integrity and Privacy)

- ▶ Joint work with Iddo Ben-Tov, Yinon Horesh and Michael Riabzev
- ▶ work in progress
- ▶ ZK, universal, succinct, scalable, transparent, post-quantum secure

A pair of novel transparent CIP

1 SCI (Scalable Computational Integrity)

- ▶ Joint work with Iddo Ben-Tov, Alessandro Chiesa, Ariel Gabizon, Daniel Genkin, Matan Hamilis, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Eran Tromer and Madars Virza
- ▶ To appear in Eurocrypt 2017
- ▶ universal, succinct, scalable, transparent, post-quantum secure

2 SCIP (Scalable Computational Integrity and Privacy)

- ▶ Joint work with Iddo Ben-Tov, Yinon Horesh and Michael Riabzev
- ▶ work in progress
- ▶ ZK, universal, succinct, scalable, transparent, post-quantum secure

Given popularity of SNARKs ...

A pair of novel transparent CIP

1 SCI (Scalable Computational Integrity)

- ▶ Joint work with Iddo Ben-Tov, Alessandro Chiesa, Ariel Gabizon, Daniel Genkin, Matan Hamilis, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Eran Tromer and Madars Virza
- ▶ To appear in Eurocrypt 2017
- ▶ universal, succinct, scalable, transparent, post-quantum secure

2 STARK (Succinct Transparent ARgument of Knowledge)

- ▶ Joint work with Iddo Ben-Tov, Yinon Horesh and Michael Riabzev
- ▶ work in progress
- ▶ ZK, universal, succinct, scalable, transparent, post-quantum secure

Given popularity of SNARKs ...



Comparison with prior CIP implementations

- Linear PCP (LPCP) [IK007]
 - ▶ Use additively homomorphic encryption to (i) hide queries + (ii) eliminate need for low-degree testing
 - ▶ Implementations: pepper, ginger, . . . [SBW11,SVP+12,SMBW12]

Comparison with prior CIP implementations

- Linear PCP (LPCP) [IK007]
- MPC-in-head (MPCh) [IKOS07]
 - ▶ Prover commits to MPC transcript, then opens one party's view
 - ▶ Implementation: ZKBoo [GM016]

Comparison with prior CIP implementations

- Linear PCP (LPCP) [IK007]
- MPC-in-head (MPCh) [IKOS07]
- Proofs for muggles (IP) [GKR08]
 - ▶ Scaling-down of $IP=PSPACE$ to case of poly-bounded prover, works for uniform log-space PTIME
 - ▶ Implementations: [CTY11,CMT12,T13], allspice [vsBW13]

Comparison with prior CIP implementations

- Linear PCP (LPCP) [IKO07]
- MPC-in-head (MPCh) [IKOS07]
- Proofs for muggles (IP) [GKR08]
- Pairing-based/Knowledge of Exponent (KOE) [G10,GGP10,L12,GGPR13]
 - ▶ succinct proofs (< 300 bytes) after setup, which is non-Arthur-Merlin
 - ▶ Implementations: Pinocchio [PGHR13], SNARKs for C [BCGTV13], Zaatar [SBVBPW13], Buffet [WSHRBW15]

Comparison with prior CIP implementations

- Linear PCP (LPCP) [IKO07]
- MPC-in-head (MPCh) [IKOS07]
- Proofs for muggles (IP) [GKR08]
- Pairing-based/Knowledge of Exponent (KOE) [G10,GGP10,L12,GGPR13]
- Discrete-logarithm problem (DLP) based [G11,S11]
 - ▶ succinct proofs, public (Arthur-Merlin) setup, verification-time $> \mathcal{T}$
 - ▶ Implementation: [BCCGP16]

Comparison with prior CIP implementations

- Linear PCP (LPCP) [IKO07]
- MPC-in-head (MPCh) [IKOS07]
- Proofs for muggles (IP) [GKR08]
- Pairing-based/Knowledge of Exponent (KOE) [G10,GGP10,L12,GGPR13]
- Discrete-logarithm problem (DLP) based [G11,S11]
- Incrementally verifiable computation (IVC) [V08,BCCT13]
 - ▶ Prover runs verifier on each prior “chunk” of computation
 - ▶ Implementation: [BCTV14] (KOE based)

Comparison of implemented CIP

- **UN universal:** works for any language in NP
- **SC scalable:** prover runtime quasilinear in \mathcal{T}
- **NI noninteractive:** after setup/common reference string
- **SU succinct:** verifier time $\text{poly}(\log \mathcal{T}, |x|)$
- **TR transparent:** setup+queries are merely public random coins
- **ZK:** zero knowledge
- **PQ:** post-quantum resistant

| | UN | SC | NI | SU | TR | ZK | PQ |
|---------------|----|----|----|----|----|----|----|
| LPCP [IK007] | + | - | - | ± | - | + | - |
| MPCh [IKOS07] | + | - | + | - | + | + | + |
| IP [GKR08] | - | - | - | + | + | - | + |
| KOE [GGPR13] | + | + | + | ± | - | + | - |
| DLP [BCCGP16] | + | + | + | - | + | + | - |
| IVC [V08] | + | + | + | + | - | + | - |

Comparison of implemented CIP

- **UN universal:** works for any language in NP
- **SC scalable:** prover runtime quasilinear in \mathcal{T}
- **NI noninteractive:** after setup/common reference string
- **SU succinct:** verifier time $\text{poly}(\log \mathcal{T}, |x|)$
- **TR transparent:** setup+queries are merely public random coins
- **ZK:** zero knowledge
- **PQ:** post-quantum resistant

| | UN | SC | NI | SU | TR | ZK | PQ |
|---------------|----|----|----|----|----|----|----|
| LPCP [IKO07] | + | - | - | ± | - | + | - |
| MPCh [IKOS07] | + | - | + | - | + | + | + |
| IP [GKR08] | - | - | - | + | + | - | + |
| KOE [GGPR13] | + | + | + | ± | - | + | - |
| DLP [BCCGP16] | + | + | + | - | + | + | - |
| IVC [v08] | + | + | + | + | - | + | - |
| SCI [BBC+17] | + | + | + | + | + | - | + |

Comparison of implemented CIP

- **UN universal:** works for any language in NP
- **SC scalable:** prover runtime quasilinear in \mathcal{T}
- **NI noninteractive:** after setup/common reference string
- **SU succinct:** verifier time $\text{poly}(\log \mathcal{T}, |x|)$
- **TR transparent:** setup+queries are merely public random coins
- **ZK:** zero knowledge
- **PQ:** post-quantum resistant

| | UN | SC | NI | SU | TR | ZK | PQ |
|----------------|----|----|----|-------|----|----|----|
| LPCP [IK007] | + | - | - | \pm | - | + | - |
| MPCh [IKOS07] | + | - | + | - | + | + | + |
| IP [GKR08] | - | - | - | + | + | - | + |
| KOE [GGPR13] | + | + | + | \pm | - | + | - |
| DLP [BCCGP16] | + | + | + | - | + | + | - |
| IVC [V08] | + | + | + | + | - | + | - |
| SCI [BBC+17] | + | + | + | + | + | - | + |
| STARK [BBHT17] | + | + | + | + | + | + | + |

SCI vs. other CIP implementations

Table: Execution of same TinyRAM program for 2^{16} cycles; 80-bit security level; machine w/ 32 AMD Opteron cores, clock rate 3.2 GHz, 512 GB RAM.

| | | KOE | IVC | DLP | SCI | STARK |
|-------------------|------------------|-----------|-----------|-----------|-----------|-----------|
| Ver. setup | time | ~ 28 min | ~ 10 sec | ~ 0.7 sec | <0.01 sec | <0.01 sec |
| | key len | ~ 18.9 GB | 43 MB | 154 MB | 16 bytes | 16 bytes |
| Prov | time | ~ 18 min | 4.2 days | ~ 8 min | ~ 41 min | 6.7 min |
| | memory | ~ 216 GB | 2.9 GB | ~ 1 TB | ~ 135 GB | ~ 131 GB |
| Ver. dec. | time | < 10 ms | ~ 25 ms | ~ 1.7 min | ~ 0.5 sec | ~ 0.1 sec |
| | comm comp | 230 bytes | 374 bytes | 8.8KB | ~ 42.5 MB | 1.8 MB |
| V. total | time | ~ 28 min | ~ 10 sec | 1.7 min | ~ 0.5 sec | ~ 0.1 sec |
| | comm comp | ~ 18.9 GB | 43 MB | ~ 154 MB | ~ 42.5 MB | ~ 1.8 MB |

Overview

- Computational integrity and privacy (CIP) — motivation ✓
- Importance of transparency ✓
- A pair of new transparent CI(P) systems
 - ▶ SCI performance [BBCGGHPRSTV16]
 - ▶ STARK performance [BBHT17]

SCI — Succinct Computational Integrity

- First assembly-code-to-PCP* reduction, including RS-proximity testing and PCPP composition

SCI — Succinct Computational Integrity

- First assembly-code-to-PCP* reduction, including RS-proximity testing and PCPP composition
(* Use IOPP instead of PCPP to save space, similar “code complexity” /security/soundness)

SCI — Succinct Computational Integrity

- First assembly-code-to-PCP* reduction, including RS-proximity testing and PCPP composition
(* Use IOPP instead of PCPP to save space, similar “code complexity”/security/soundness)
- Joint work with Iddo Ben-Tov, Alessandro Chiesa, Ariel Gabizon, Daniel Genkin, Matan Haimis, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Eran Tromer and Madars Virza [Eurocrypt17]

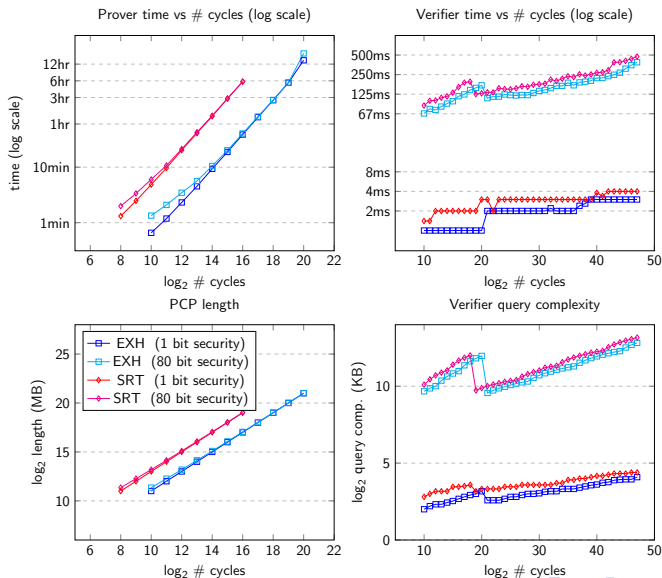
SCI — Succinct Computational Integrity

- First assembly-code-to-PCP* reduction, including RS-proximity testing and PCPP composition
(* Use IOPP instead of PCPP to save space, similar “code complexity” /security/soundness)
- Joint work with Iddo Ben-Tov, Alessandro Chiesa, Ariel Gabizon, Daniel Genkin, Matan Haimis, Evgenya Pergament, Michael Riabzev, Mark Silberstein, Eran Tromer and Madars Virza [Eurocrypt17]
- Large scale effort
 - ▶ Started summer 2010
 - ▶ More than 1M euro over first 6 years (thanks to European Research Council!!)
 - ▶ Mostly for programmers: Ohad Barta, Lior Greenblatt, Shaul Kfir, Gil Timnat, Arnon Yogev

SCI executed programs

- CI statement: “no subset of input array A sums to target t ”
- Two different programs
 - 1 EXH — exhaustive search
 - ★ running time $\mathcal{T} \sim 2^{|A|}$
 - ★ memory $O(1)$
 - 2 SRT — sorted search
 - ★ Sort each half of A increasingly, then “merge”
 - ★ running time $\mathcal{T} \sim 2^{|A|/2}$
 - ★ random access memory consumption $2^{|A|/2}$


SCI numbers




SCI break-even point [SVPBBW12,SMBW12]

- Def: minimal n_0 for which naïve re-execution $>$ SCI-verification.
- For EXH at 80-bit security $n_{\text{EXH}} = 22$
- For SRT at 80-bit security $n_{\text{SRT}} = 48$
- $n_{\text{SRT}} > n_{\text{EXH}}$ because SRT is quadratically faster

Cutoff point for EXH



Cutoff point for SRT



STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;

STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;
- also first implemented post-quantum secure succinct zk proof system

STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;
- also first implemented post-quantum secure succinct zk proof system
- uses recent scalable “duplex IOP” BCGV16, along with
 - 1 new “biased RS-IOPP” protocol

STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;
- also first implemented post-quantum secure succinct zk proof system
- uses recent scalable “duplex IOP” BCGV16, along with
 - 1 new “biased RS-IOPP” protocol
 - ★ better concrete soundness and security than prior state of the art
 - ★ shorter proofs due to better “packaging” of proof parts into a Merkle tree
 - ★ more efficient to compute (single FFT followed by fully parallelizable local computations)

STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;
- also first implemented post-quantum secure succinct zk proof system
- uses recent scalable “duplex IOP” BCGV16, along with
 - ① new “biased RS-IOPP” protocol
 - ★ better concrete soundness and security than prior state of the art
 - ★ shorter proofs due to better “packaging” of proof parts into a Merkle tree
 - ★ more efficient to compute (single FFT followed by fully parallelizable local computations)
 - ② new IOP reduction from Algebraic constraint satisfaction to Reed-Solomon proximity testing
 - ★ higher soundness retention
 - ★ simpler proofs, of lower-degree

STARK [BBHT17]

- first implementation of a succinct, universal, transparent, scalable zk proof system;
- also first implemented post-quantum secure succinct zk proof system
- uses recent scalable “duplex IOP” BCGV16, along with
 - ① new “biased RS-IOPP” protocol
 - ★ better concrete soundness and security than prior state of the art
 - ★ shorter proofs due to better “packaging” of proof parts into a Merkle tree
 - ★ more efficient to compute (single FFT followed by fully parallelizable local computations)
 - ② new IOP reduction from Algebraic constraint satisfaction to Reed-Solomon proximity testing
 - ★ higher soundness retention
 - ★ simpler proofs, of lower-degree
 - ③ improved concrete arithmetization of cryptographic primitives (AES)

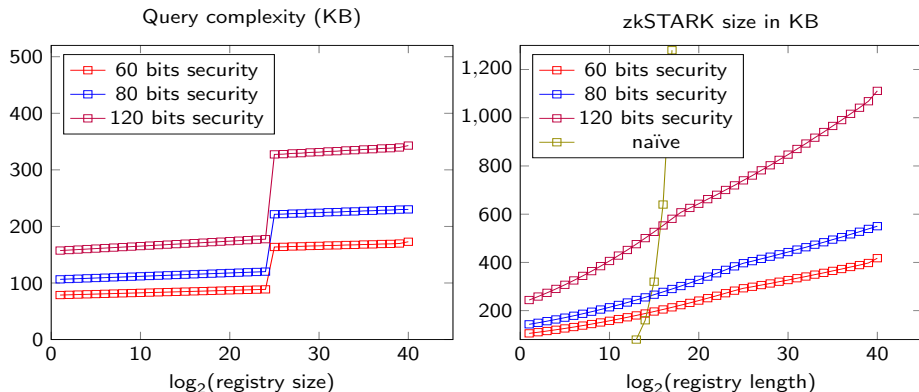
STARK black-list non-membership [BBHT17]

- CIP statement: y does not appear in private black-list, with public hash commitment r
- Useful for banks (FACTA compliance), airlines (anti-terrorism compliance), etc.

STARK black-list non-membership [BBHT17]

- CIP statement: y does not appear in private black-list, with public hash commitment r
- Useful for banks (FACTA compliance), airlines (anti-terrorism compliance), etc.
- Formally
 - ▶ Inputs: element y and public commitment r (hash of black-list)
 - ▶ Statement: $\exists D \text{ comm}(D) = r \text{ and } y \notin D$
 - ▶ pseudo-code: If $y \in D$ or $\text{comm}(D) \neq r$ reject, else accept
 - ▶ D is private (nondeterministic witness), $|D| = 2^h$
 - ▶ comm is either Merkle tree or hash chain
 - ▶ Hash function is Davies-Meyer hash + AES160

STARK estimated proof length ^[BBHT17]



Disclaimer: work in progress, hence numbers may change

SCI vs. other CIP implementations

Table: Execution of same TinyRAM program for 2^{16} cycles; 80-bit security level; machine w/ 32 AMD Opteron cores, clock rate 3.2 GHz, 512 GB RAM.

| | | KOE | IVC | DLP | SCI | STARK |
|------------|-----------|-----------|-----------|-----------|-----------|-----------|
| Ver. setup | time | ~ 28 min | ~ 10 sec | ~ 0.7 sec | <0.01 sec | <0.01 sec |
| | key len | ~ 18.9 GB | 43 MB | 154 MB | 16 bytes | 16 bytes |
| Prov | time | ~ 18 min | 4.2 days | ~ 8 min | ~ 41 min | 6.7 min |
| | memory | ~ 216 GB | 2.9 GB | ~ 1 TB | ~ 135 GB | ~ 131 GB |
| Ver. dec. | time | < 10 ms | ~ 25 ms | ~ 1.7 min | ~ 0.5 sec | ~ 0.1 sec |
| | comm comp | 230 bytes | 374 bytes | 8.8KB | ~ 42.5 MB | 1.8 MB |
| V. total | time | ~ 28 min | ~ 10 sec | 1.7 min | ~ 0.5 sec | ~ 0.1 sec |
| | comm comp | ~ 18.9 GB | 43 MB | ~ 154 MB | ~ 42.5 MB | ~ 1.8 MB |

STARK vs. SNARK

- Main advantages of STARK over SNARK are transparency and scalability
 - ▶ both due to reliance on proven mathematics (PCPs) which lead to “lighter” crypto assumptions (hash+Fiat Shamir)
- Main advantage of SNARK of STARK is shorter proofs

STARK vs. SNARK

- Main advantages of STARK over SNARK are transparency and scalability
 - ▶ both due to reliance on proven mathematics (PCPs) which lead to “lighter” crypto assumptions (hash+Fiat Shamir)
- Main advantage of SNARK of STARK is shorter proofs
- Assuming STARK proofs don't get shorter, to use in a crypto-currency:
 - ▶ users send tx to a “tx-aggregator”
 - ▶ tx-aggregator checks and aggregates many transactions
 - ▶ generates single STARK for all of them (say, 2^{20})
 - ▶ broadcasts UTXO diff file + STARK
 - ▶ this improves the crypto-currency scalability
 - ▶ transparency implies: don't trust aggregator, trust the proof.

Concluding remarks

- Computational integrity+privacy (CIP)
 - ▶ crucial for long-term viability of decentralized blockchains
 - ▶ potentially useful even for trusted parties (Government, Banks, etc.)
- CIP systems for blockchains require universality, transparency, succinctness, scalability, and privacy (post-quantum security also helpful)
- STARK delivers all; SCI delivers all but privacy

Concluding remarks

- Computational integrity+privacy (CIP)
 - ▶ crucial for long-term viability of decentralized blockchains
 - ▶ potentially useful even for trusted parties (Government, Banks, etc.)
- CIP systems for blockchains require universality, transparency, succinctness, scalability, and privacy (post-quantum security also helpful)
- STARK delivers all; SCI delivers all but privacy
- “theory-to-practice” research, like this, leads to new models, new questions, and new applications

Concluding remarks

- Computational integrity+privacy (CIP)
 - ▶ crucial for long-term viability of decentralized blockchains
 - ▶ potentially useful even for trusted parties (Government, Banks, etc.)
- CIP systems for blockchains require universality, transparency, succinctness, scalability, and privacy (post-quantum security also helpful)
- STARK delivers all; SCI delivers all but privacy
- “theory-to-practice” research, like this, leads to new models, new questions, and new applications
- want to hear more?
 - ▶ Ethereum meetup this Sunday Jan 29, 6pm, Institute for the Future: *more details*
 - ▶ Berkeley CS Theory Seminar, Monday Feb 6, 4pm, Wozniak Lounge: *moon math*
 - ▶ Stanford Security Seminar, Tuesday Feb 7, 4:15pm, Gates 463: *moon math+engineering*